


Dashboards and feature gaps

Where the two stop being the same kind of product.


Dashboards



20 traces, \$0.036312 total, 11.17K tokens, 198 observations split DEFAULT 168 / ERROR 30 , and average score 0.8 .

That error split is the single most useful number on the page. 30 out of 198 observations errored, which is the flaky weather tool showing up as a metric instead of a surprise.


LangSmith calls this Monitoring, and it goes deeper:



Tabs for Traces, LLM Calls, Cost & Tokens, Tools, Run Types and Feedback Scores. Trace Count is split into Success and Error lines, and there is a dedicated Trace Error Rate chart.

Note the second tab at the top: **Alerts**. LangSmith can notify you when error rate or latency crosses a threshold.

Langfuse has Monitors for the same idea:



What each tool has that the other does not


Up to here the two have been close. This is where they stop being the same kind of product.

LangSmith is turning into a platform for building and running agents. Langfuse is an observability and evaluation tool. That is the honest summary, and it cuts both ways.

Side by side, straight from the two menus:

Feature	LangSmith	Langfuse
Tracing	yes	yes
Dashboards and charts	yes (Monitoring)	yes
Alerts on thresholds	yes	yes (Monitors)
Datasets and experiments	yes	yes
LLM-as-judge evaluators	yes	yes
Human annotation	yes (Annotation Queues)	yes
Prompt management	yes	yes
Prompt playground	yes	yes
Grouping runs	Threads	Sessions
Per-user cost tracking	no dedicated page	yes
Self-hosting	no	yes
Visual agent debugger	Studio	no
Agent hosting	Deployments	no
Code sandboxes	Sandboxes	no
LLM proxy	LLM Gateway	no
Context Hub	yes	no


Studio, the one people actually mean



Studio draws your graph (`__start__` → `model` → `tools` → `__end__`) and lets you run it step by step, inspect state at each node, and edit that state before continuing. It connects to either a LangSmith Deployment or a graph running locally through the Agent Server.

This is the biggest single gap. Langfuse can show you a graph of a run that **already happened**. Studio lets you drive the graph **while it runs**, stop it mid-loop and change the state. Those are different jobs. If you build in LangGraph, Studio is a genuine reason to keep LangSmith even if your traces live elsewhere.

Deployments



One-click agent hosting with pre-built APIs, built for long-running and bursty workloads. Note the **Upgrade required** badge: this is a paid tier, not something you get on the free plan.

Sandboxes

LangSmith

Application

All applications

Search

Home

Tracing2

Monitoring

Datasets & Experiments1

Evaluators

Annotation Queues

Prompts

Playground

Studio

Context Hub

Deployments

Sandboxes

LLM GatewayBeta

Settings

Personal

vikashra041013@gmail.c...

Personal / Sandboxes

Sandboxes

Sandboxes are available on plans with sandbox access. Upgrade to unlock sandbox environments for this workspace.

Upgrade

Isolated environments for agents that execute code. Also paid.

LLM Gateway (Beta)

The screenshot shows the LangSmith LLM Gateway interface. On the left is a sidebar with navigation options: Application, All applications, Search, Home, Tracing, Monitoring, Datasets & Experiments, Evaluators, Annotation Queues, Prompts, Playground, Studio, Context Hub, Deployments, Sandboxes, and LLM Gateway (Beta). The main content area is titled 'LangSmith LLM Gateway' and includes a 'Gateway Credits' section showing a balance of \$0.00 and a 'Bring Your Own Key' section. A code editor displays a TypeScript snippet for using OpenAI through the gateway.

```
import OpenAI from "openai";

const client = new OpenAI({
  apiKey: process.env.LANGSMITH_API_KEY,
  baseURL: "https://gateway.smith.langchain.com/v1",
});

const response = await client.responses.create({
  model: "moonshotai/Kimi-K3",
  input: "Hello!",
});
```

A proxy in front of model providers. You authenticate to LangSmith instead of holding provider keys, and it handles cost controls, model fallbacks and spend monitoring. Point your OpenAI client at `https://gateway.smith.langchain.com/v1` and it works unchanged.

For this lab that is directly relevant. The agent currently talks to Azure AI Foundry using a key in `.env`. The Gateway is the alternative to managing that key yourself.

What Langfuse has that LangSmith does not

Self-hosting is the headline and the reason this lab runs Langfuse at all. Your traces, prompts and eval results never leave your network. For anything under a data residency rule, that single fact outweighs the entire feature list above.

Per-user cost tracking gets its own page:

langfuse v4.6.0 OSS

llm-observability-lab

llm-lab

1d Past 1 day

Go to... ctrl K

Home

Dashboards

Observability

Tracing

Sessions

Users

Monitors

Prompt Management

Prompts

Playground

Evaluation

Scores

Evaluators

Human Annotation

Datasets

Experiments

Star Langfuse

See the latest releases and help grow the community on GitHub

Langfuse 33k

Settings

Book a call

Support

A admin

admin@jmagro...

Users

Search (User ID)

Env default sdk-experiment

Filters

User ID	Environment	First Event	Last Event	Total Events	Total Tokens	Total Cost
vikash	default	8/7/2026, 1:08:53 AM	8/7/2026, 12:43:13 PM	102	6.36K	\$0.02

Rows per page 50

Page 1 of 1

<<

<

>

>>

User `vikash` , 102 events, 6.36K tokens, \$0.02 spent. LangSmith stores `user_id` as ordinary metadata you can filter on. Langfuse aggregates it into a real screen. If you need to know what a specific customer costs you, Langfuse answers it in one click.

Evaluators run against live production traces, configured in the UI:

langfuse

v4.6.0 OSS

llm-observability-lab / llm-lab

Go to...

ctrl K

Home

Dashboards

Observability

Tracing

Sessions

Users

Monitors

Prompt Management

Prompts

Playground

Evaluation

Scores

Evaluators

Human Annotation

Datasets

Experiments

Star Langfuse

See the latest releases and help grow the community on GitHub

Langfuse 33k

Settings

Book a call

Support

admin admin@jmgro...

Evaluators

Get Started with LLM-as-a-Judge Evaluations

Create evaluation templates and evaluators to automatically score your traces with LLM-as-a-judge. Set up custom evaluation criteria and let AI help you measure the quality of your outputs.

Create Evaluator

Learn More

Create eval template

Context relevance

name

context-relevance

Prompt

Example: George has visited service before your assistant. Customer describes his needs. "Get your service, he has great insights!" So there are notes to this statement or do it a bunch of history? Well no, service won't answer your stated needs if you have less than perfect vision. A list of service won't give a blind person their vision. If your vision problem isn't related to vision, it, your vision won't change no matter how many services you use. Service is a... However! The first sentence is introducing the topic of the query but not relevant to answer it. The following statement clearly answers the question and thus, is relevant. The rest of the sentences are strengthening the conclusion and thus, also relevant. Output: George (assistant) (context: Relevant)[["Think step by step."]]

Model

Provider

Model name

Temperature

Top P

API key

Score

provide a score between 0 and 1

You use function calls to extract data from the LLM. Specify what the LLM should return for the score.

Automate evaluations

Use LLM-as-a-judge to automatically evaluate your traces without manual review

Measure quality

Create custom evaluation criteria to measure the quality of your LLM outputs

Scale efficiently


Evaluate thousands of traces automatically with customizable sampling rates

Track performance


Monitor evaluation metrics over time to identify trends and improvements

You build the judge template in a form (prompt, model, temperature, sampling rate) and it scores real traffic automatically. No code, no dataset required. LangSmith has evaluators too, but the Langfuse version is aimed squarely at continuous scoring of production traces rather than at offline dataset runs.


Prompt management and Playground exist on both sides, so neither wins here:



Langfuse Playground:



LangSmith Playground:



How to choose

- Building on LangGraph and you want to debug agents visually, host them, or proxy your models? **LangSmith**, and it is not close. Studio alone decides it.
- Data cannot leave your network, or you need per-user cost attribution, or you want judges scoring live traffic without writing code? **Langfuse**.
- Just want to see what your agent did and score it against a dataset? Either. They are close enough that the deciding factor is hosting, not features.